

Proof-Guided Dereference Collapse

Loop-Invariant Hoisting of Guarded Pointer Chains in LLVM

January 28, 2026

- ▶ C kernels use checked pointer operations (no undefined behavior, UB)
- ▶ LLVM pass builds a guarded heap-access graph (proof of purity)
- ▶ Collapse pass hoists invariant deref chains into a preheader (SSA)

1. Motivation

- ▶ Pointer dereference chains inside loops are expensive and branchy
- ▶ Guard checks make safety explicit but add repeated work
- ▶ If inputs are invariant, repeating guarded chains is pure waste

Goal

Replace repeated guarded dereferences in a loop with a single cached SSA value.

Scope: Checked-pointer API

- ▶ The LLVM pass only recognizes explicit `ck_*` calls
- ▶ Raw C pointer ops are not analyzed or rewritten
- ▶ Safety is encoded in `ck_*` guards + loads (no UB)
- ▶ Whole-program use would require a separate source-to-source rewrite

2. Baseline C code

Baseline

```
for (uint64_t k = 0; k < iters; ++k) {  
    Eval r = triple_deref(heap, p, 0);  
    sum += r.value;  
}
```

- ▶ Each iteration performs guards + 3 loads
- ▶ Work scales with loop trip count
- ▶ No reuse across iterations

Analysis: Baseline C

- ▶ Each iteration re-enters the guarded dereference chain.
- ▶ All guards and loads are repeated even when inputs are invariant.
- ▶ The loop body cost scales with chain length \times trip count.

3. Optimized C (conceptual)

Hoisted

```
Eval cached = triple_deref(heap, p, 0);
for (uint64_t k = 0; k < iters; ++k) {
    Eval r = cached;
    sum += r.value;
}
```

- ▶ Compute once in preheader
- ▶ Loop uses SSA value only
- ▶ Eliminates repeated guards + loads

Analysis: Hoisted C

- ▶ The guarded chain is evaluated once in the preheader.
- ▶ The loop consumes a single SSA value per iteration.
- ▶ Repeated guards and loads are removed from the hot path.

4. Guarded semantics (baseline vs optimized)

Baseline (inside loop)

```
vp  = guard_ptr(p)
v1  = guard_nonnull(vp)
v2  = load_ptr(v1)
v3  = guard_ptr(v2)
v4  = guard_nonnull(v3)
v5  = load_ptr(v4)
v6  = guard_ptr(v5)
v7  = guard_nonnull(v6)
v8  = load_ptr(v7)
return v8
```

Optimized

```
// preheader
vp  = guard_ptr(p)
...
cached = load_ptr(v7)

// inside loop
return cached
```

Analysis: Guarded Semantics

- ▶ Baseline repeats a pure chain of guards + loads.
- ▶ Optimized version preserves the same checks but runs them once.
- ▶ Semantic proof: the guarded chain is side-effect free.

5. LLVM IR (baseline)

```
loop:  
  %val = call { i64, i32 } @triple_deref(...)  
  %v    = extractvalue { i64, i32 } %val, 1  
  %sum = add i64 %sum, %v  
  br label %loop
```

- ▶ call @triple_deref sits inside the loop body
- ▶ Repeated guard+load chain every iteration

Analysis: LLVM IR Baseline

- ▶ The call is inside the loop, so costs multiply by trip count.
- ▶ The loop depends on a pointer-chasing call each iteration.
- ▶ No reuse of the computed value.

6. LLVM IR (optimized, SSA)

```
preheader:  
  %hoisted = call { i64, i32 } @triple_deref(...)  
  br label %loop  
  
loop:  
  %v = extractvalue { i64, i32 } %hoisted, 1  
  %sum = add i64 %sum, %v  
  br label %loop
```

- ▶ One call in preheader, no loads from cache
- ▶ SSA value dominates all loop uses

Analysis: LLVM IR Optimized

- ▶ The call is hoisted to the preheader and dominates the loop.
- ▶ The loop uses a single SSA value – no memory loads.
- ▶ This is loop-invariant code motion with a proof of purity.

7. Assembly intuition

Baseline

```
loop:  
    call triple_deref ; guard+load chain  
    add s0, s0, a0  
    j     loop
```

Optimized

```
preheader:  
    call triple_deref  
loop:  
    add s0, s0, a0      ; cached SSA value  
    j     loop
```

Analysis: Assembly Intuition

- ▶ Baseline spends cycles in a call + dependent loads per iteration.
- ▶ Optimized loop is reduced to arithmetic on a cached value.
- ▶ Pointer chasing is removed from the hot loop body.

8. Correctness argument

- ▶ **Loop invariance:** inputs to `triple_deref` are invariant in the loop.
- ▶ **Purity:** guarded deref graph contains only guards + loads, no side effects.
- ▶ **Dominance:** preheader dominates loop body, so SSA value is available.

Key reasoning

If f is pure and x is loop-invariant, then $f(x)$ is loop-invariant. Hoisting preserves semantics.

9. Benchmark methodology

- ▶ **Unoptimized benchmark:** volatile sink inside loop (noisy, hides benefit).
- ▶ **SSA benchmark:** register accumulator + single compiler barrier.
- ▶ Same heap layout and `triple_deref` implementation for comparability.
- ▶ Run config: 50,000,000 iterations, 7 runs.

10. Benchmark results

Variant	Base mean (ms)	Base ns/iter	Opt mean (ms)	Opt ns/iter	Speedup
Unoptimized	1632.35	32.647	95.75	1.915	17.30
SSA	1642.58	32.852	61.39	1.228	26.88

Table: Mean timings over 7 runs at 50,000,000 iterations.

- ▶ Microbenchmarks amplify savings (loop body becomes very small)
- ▶ SSA benchmark isolates the optimization effect and lowers variance

11. Generalization

- ▶ Extend beyond `triple_deref` to any linear guarded dereference chain
- ▶ Recognize patterns at IR level (guards + loads + fields)
- ▶ Hoist when alias + memory stability are proven (MemorySSA/AA)

12. From demo to production optimizer

- ▶ **Production-ready:** SSA hoist, dominance checks, preheader insertion
- ▶ **Missing for real C/C++:** MemorySSA + alias analysis for heap stability
- ▶ **Volatile/atomic handling:** must avoid speculation across side effects
- ▶ **Profitability:** TTI cost models, loop trip counts, reg pressure
- ▶ **Robustness:** remove dependence on runtime JSON graphs
- ▶ **Still valid:** proof-guided LICM at IR level, target-independent